CASE STUDY

TaliScan

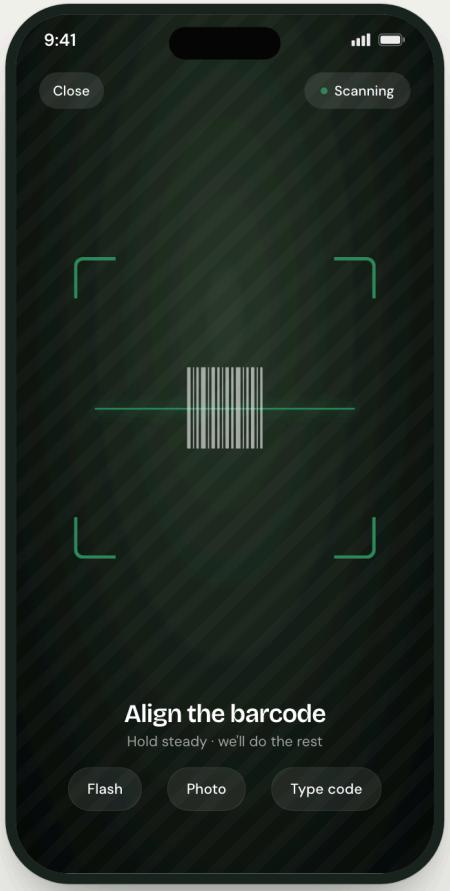
An iPhone app that shows families what's really inside their food, personal-care, and cleaning products, and scores each one on the science, not the marketing.

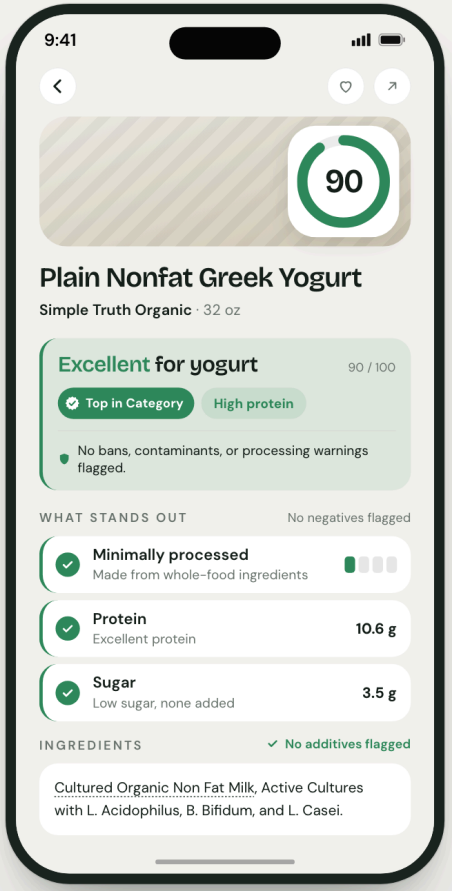
- Built solo
- Launched on the App Store, May 2026

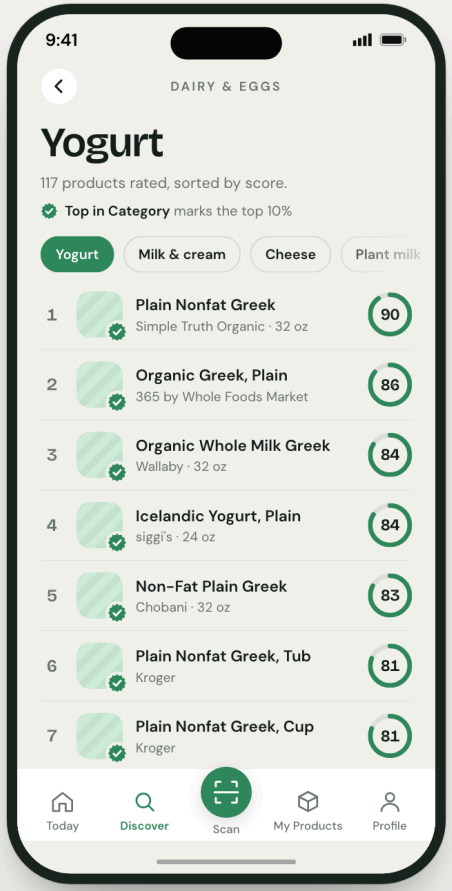
Rick Gibb

Founder & Product Lead

github.com/gibbitall/taliscan-case-study



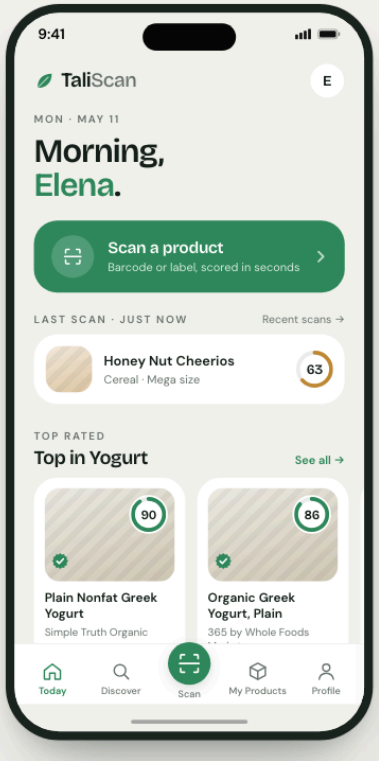




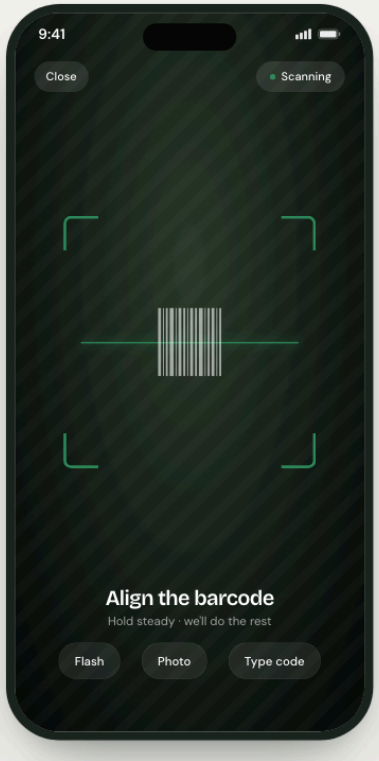
THE PRODUCT

One scan, a clear answer.

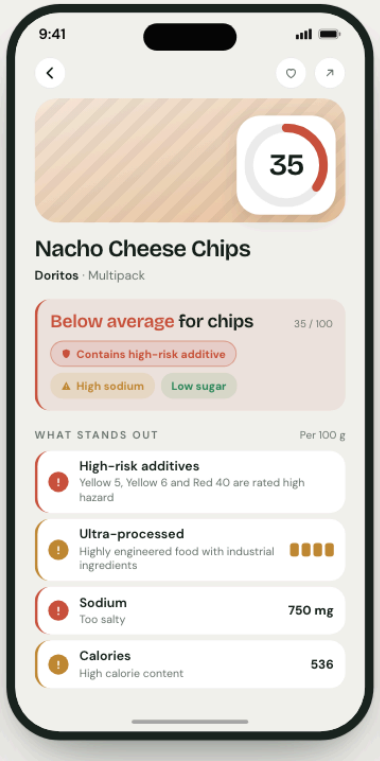
One 0–100 scale for food, personal care, and cleaning products. Beyond the score: tap-to-learn ingredients, top lists by category, and a home audit that shows when a saved product's score changes.



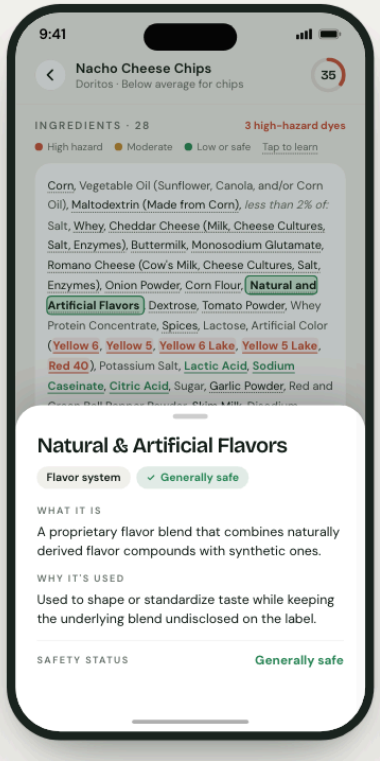
1 Home



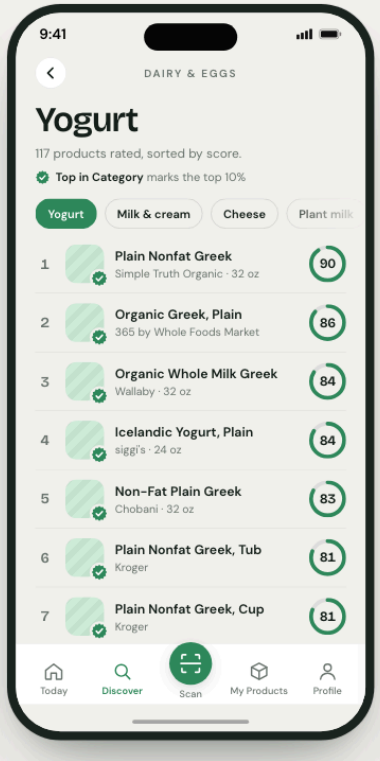
2 Scan



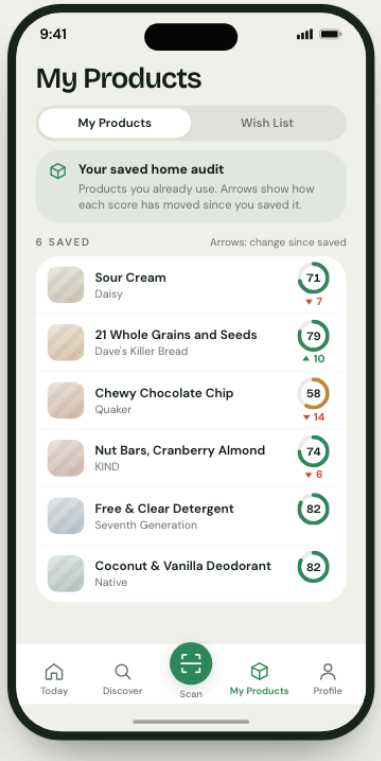
3 Score and flags



4 Tap to learn



5 Top in Category

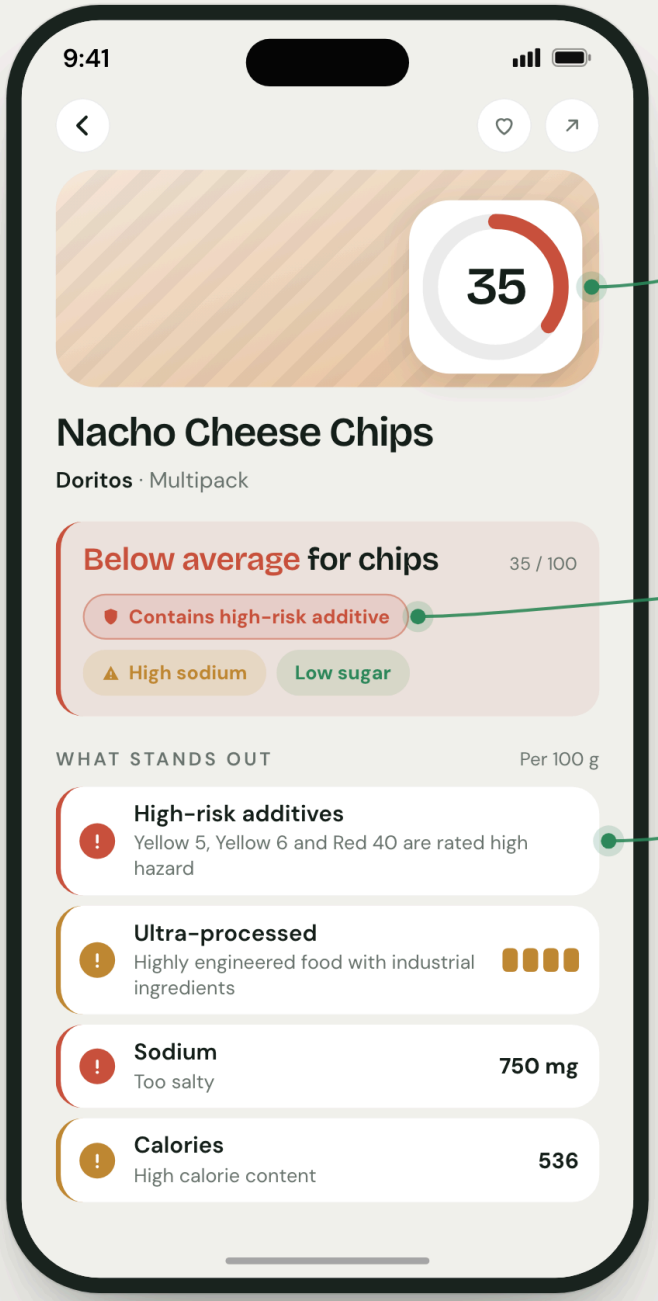


6 Home audit

THE KEY DECISION

Every score shows its work.

Health guidance should give the same answer every time and explain itself, so the score is a deterministic rules engine. AI does the slow work that scales: researching the evidence behind each rule and building the product catalog.



1 A 0 to 100 score, in context

Below average for chips. A rules engine computes it with no LLM at runtime, so the same product always gets the same score.

2 Safety flags

Yellow 5, Yellow 6 and Red 40 are rated high hazard. Two or more high-hazard additives cap a score at 35.

3 Every deduction explained

Ultra-processing, sodium and calories each show the reason, with nutrient amounts per 100 g.

HOW IT WORKS

AI where it scales. Rules where trust matters.

AI agents do the slow research ahead of time, in batch. When someone scans, only rules produce the score.

AI · Rules in code · My call

AHEAD OF TIME · AI, IN BATCH

Products to add
Retailer listings, plus barcodes shoppers scanned that weren't in the catalog

→

AI research
Retailer, manufacturer, and public sources, reconciled into one record

→

Quality gate
8 core fields, checked in code before anything goes live

→

Live catalog
Nearly 7,000 products at launch, across 21 categories

Agent team Records that fail the gate go to a lead agent that runs the loop and six specialists: a researcher, variant linker, read-only validator, senior reviewer, reject-by-default gatekeeper, and monitor. They repair records overnight and send them back through the gate. Hard cases come to me.

Published evidence
2023–2026 PubMed, EFSA, FDA, and IARC sources

→

AI research agents
Evidence notes and a strength grade for each of 307 additives

→

My call
Which of the 17 proposed rating changes to adopt

→

Hazard ratings
Weighted by how each ingredient reaches the body

AT SCAN TIME · RULES ONLY

Scan a barcode
On the phone

→

Rules engine
Live catalog plus hazard ratings. No LLM.

→

Score in context
0–100, with every reason shown

New barcode?
Label photos give a provisional score, and the product joins the research queue

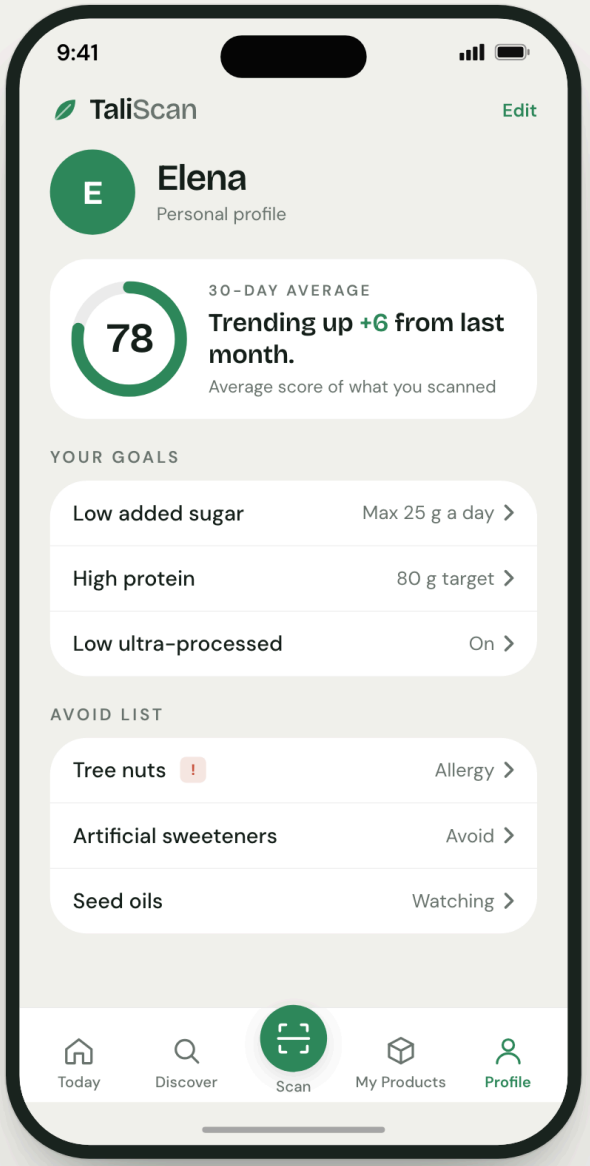
WHERE IT WAS GOING

Next: a verdict made for you.

DESIGN CONCEPT

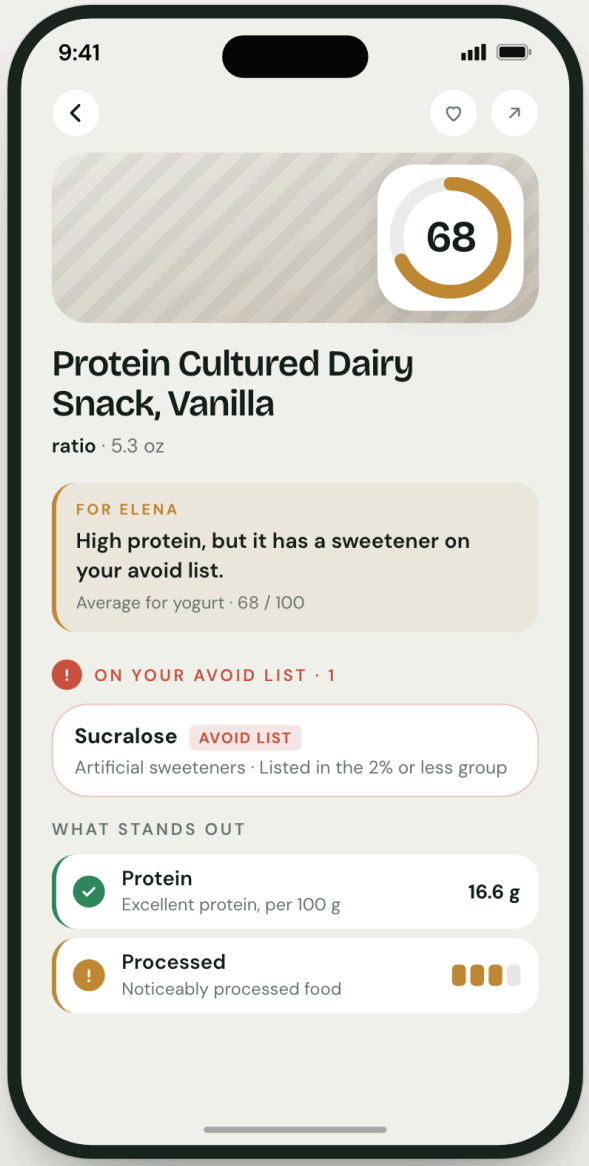
Version 1.0 scored each product the same way for everyone. The next step was personal: goals and an avoid list, with a "for you" verdict layered over the shared score instead of replacing it.

Product, score, and ingredients are real catalog data.



CONCEPT

Goals and an avoid list



CONCEPT

A verdict for one person